

Software Development ***Python (Part C)***



Davide Balzarotti

Eurecom – Sophia Antipolis, France

Overview

- Modules and packages
- Python standard library
 - The sys module
 - The os (Operating System) module
 - Spawning and controlling other processes
 - Regular expressions
 - Network
 - Threads
 - Signals

Modules

- Modules provide a basic way of organizing code
- In Python, a module is simply a Python file
 - Any python script can be seen (and imported) as a module
 - Module X is just a file named X.py
- Modules are imported using the `import` keyword
- When you write `import foo` Python:
 - If the module `foo` has already been imported, it does nothing (it is not imported again!)
 - Otherwise, it searches for the file `foo.py` (or `foo.pyc`) in the search path
 - By default, the search path includes the current working directory and the standard library directories

Modules

- Each module has a symbol table (`__dict__`) that contains all the names defined by the module
- A module usually contains variables and function (or class) definitions. But it can also contain normal code
 - When a module is imported, all the instructions contained in the python file are executed
 - Within a module, the module's name is available as the value of the global variable `__name__`
 - If you want to execute some code only if the file is executed as a script (not when imported as a module) you can use:

```
if __name__ == "__main__":
```
- `reload(modulename)` allows to re-import a module that has already been imported

Modules

```
# foo.py

def f(a,b,c):
    return a+b+c

print 'Hi'
if __name__ == '__main__':
    print 'I was directly invoked'
else:
    print 'I was imported as a module'
```

```
>> import foo
```

```
Hi
```

```
I was imported as a module
```

```
>> foo.f(1,2,3)
```

```
6
```

Modules

```
# foo.py

def f(a,b,c):
    return a+b+c

print 'Hi'
if __name__ == '__main__':
    print 'I was directly invoked'
else:
    print 'I was imported as a module'
```

```
>> import foo
Hi
I was imported as a module
>> foo.f(1,2,3)
6
```

```
>> from foo import f
Hi
I was imported as a module
>> import foo
>> reload(foo)
Hi
I was imported as a module
<module 'foo' from 'foo.pyc'>
```

Modules

```
# foo.py

def f(a,b,c):
    return a+b+c

print 'Hi'
if __name__ == '__main__':
    print 'I was directly invoked'
else:
    print 'I was imported as a module'
```

```
>> import foo
Hi
I was imported as a module
>> foo.f(1,2,3)
6
```

```
>> from foo import f
Hi
I was imported as a module
```

```
>> import foo
```

```
>> reload(foo)
```

```
Hi
```

```
I was imported as a m
<module 'foo' from 'fo
```

```
balzarot> python foo.py
```

```
Hi
```

```
I was directly invoked
```

```
balzarot>
```

Packages

- Packages are a way to group together and organize a set of related modules
- Module names in a package are accessible using the *dot* notation
 - For example, the module name A.B designates a module named B in a package named A
- The structure of the package hierarchy is determined by the organization of the modules in the filesystem
 - Useful since putting many modules in the same directory can be cumbersome

Packages

- A package in Python is simply a directory that contains a file named `__init__.py`
 - The `__init__.py` file can execute some initialization code, or it can be an empty file if no initialization is required

```
sound/                Top-level package
  __init__.py         Initialize the sound package
  formats/           File format subpackage
    __init__.py
    wavread.py
    wavwrite.py
    aiffread.py
    aiffwrite.py
    auwrite.py
    ...
  effects/           Subpackage for sound effects
    __init__.py
    echo.py
    surround.py
    reverse.py
    ...
```

```
>>> import sound.effects.echo
>>> sound.effects.echo.echofilter(input, output, delay=0.7)
```

Importing * From a Package

- In presence of complex packages, `import *` could take a long time and have unwanted side-effects
- By default, when `import *` is used on a package it does not import all submodules into the current namespace
- The only solution is for the package author to provide an explicit index of the package
 - If a package's `__init__.py` code defines a list named `__all__`, it is taken to be the list of module names that should be imported when the user use `import *`

Today

- Modules and packages
- Python standard library
 - The sys module
 - The os (Operating System) module
 - Spawning and controlling other processes
 - Regular expressions
 - Network
 - Threads
 - Signals

sys Module

- `path` – list of the paths where Python is looking for when it imports modules
- `version`, `subversion`, `version_info` – info about the version of Python running
- `stdin`, `stdout`, `stderr` – no need to explain
- `argv` – list of the command-line arguments given to the program
 - Each argument is always a string
 - The first argument is the name of the program
- `exit(value)` – terminates the program

Simple Parameters Handling

```
import sys

if len(sys.argv) < 2:
    sys.stderr.write("Use: %s value\n"%sys.argv[0])
    sys.stderr.write("  value - just a number\n")
    sys.exit(1)

value = int(sys.argv[1])
...
```

Getopt (the C way)

Provide an easier way to parse the command line arguments

```
import getopt
# example of parameter list (you should use sys.argv)
arglist = ['-ab', '-cfoo', '-d', 'bar', 'a1', 'a2']
# parsing the options
opts, args = getopt.getopt(arglist, 'abc:d:')

# opts=[('-a', ''), ('-b', ''), ('-c', 'foo'), ('-d', 'bar')]
# args=['a1', 'a2']
```



Short options

the ':' characters means that the option requires an argument

Getopt (the C way)

Provide an easier way to parse the command line arguments

```
import getopt
# parameters (containing long options too)
arglist = ['--cond=foo', '--testing', '--ofile', 'abc.def',
           '-x', 'a1', 'a2']
# parsing the parameters
optlist, args = getopt.getopt(args, 'x',
                               ['cond=', 'ofile=', 'testing'])

# optlist = [('--cond', 'foo'), ('--testing', ''),
             ('--ofile', 'abc.def'), ('-x', '')]
# args=['a1', 'a2']
```

Long options

the '=' characters means that the option requires an argument

ArgParse (the Python 2.7 way)

- Flexible, and powerful library for parsing command-line options
- Replace OptParse introduced (and already deprecated) in 2.6

```
from argparse import ArgumentParser
# create a new option parser
parser = ArgumentParser(description="process some integers")

# add the allowed arguments
parser.add_argument("integers", metavar="N", type=int,
                    nargs='+', help="value for the accumulator")

parser.add_argument("--sum", dest="operator", default="max",
                    action="store_const", const="sum",
                    help="sum the values (instead of finding
                        the max)")

# parse the command line arguments
args = parser.parse_args()
print "Integers:", args.integers
print "Operator:", args.operator
```


ArgParse (the Python 2.7 way)

- Flexible, and powerful library for parsing command-line options
- Replace OptParse introduced (and already deprecated) in 2.6

```
from argparse import ArgumentParser
# create a new option parser
parser = ArgumentParser(description="process some integers")
```

```
balzarot> script.py -h
# usage: script.py [-h] [--sum] N [N ...]
pa
process some integers
ad
pa positional arguments:
    N                value for the accumulator

optional arguments:
    -h, --help      show this help message and exit
#    --sum          sum the values (instead of finding the max
ar
print "Integers:", args.integers
print "Operator:", args.operator
```

OS

- The `os` module provides a portable way of using operating system dependent functionalities
- Access to the process environment
 - `environ` – dictionary containing the environment variables
- File/directory management
 - Shell-like commands
(`chdir`, `chmod`, `chown`, `mkdir`, `rmdir`, `link`..)
 - File descriptors operators: `dup`, `open`, `close`, `fstat`..
 - `listdir(path)` – return a list containing all the files and directories in *path*
- Process management:
 - `execl`, `execle`, `execv`, `popen`, `popen2`, `popen3`..

Paths Manipulation

- The `os.path` module implements some useful functions on pathnames
 - Available in different flavors to fit different operating systems (`posixpath`, `ntpath`, `macpath`..)
- Info about a file:
`getsize()`, `getatime()`, `getmtime()`, `getctime()`
- Absolute path and current directory:
`os.path.abspath(os.path.curdir)`
- Check for file/directory existence:
`exists(path)`, `isdir(path)`, `isfile(path)`
- Join multiple path together: `join(path1, path2, ...)`
- Split a path in its directory and file parts: `split(path)`

Example

```
import os
import os.path

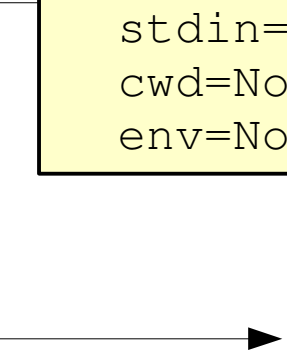
def whereis(program):
    for d in os.environ.get('PATH', '').split(':'):
        full_path = os.path.join(d, program)
        if os.path.exists(full_path) and not \
            os.path.isdir(full_path):
            return full_path
    return None
```

subprocess

- Provide functionalities to spawn new processes, connect to their input/output/error pipes, and obtain their return codes
 - To manage subprocesses, Python had many functions spread in different modules - the subprocess module intends to replace them all
- The process creation and communication is done through the `subprocess.Popen` class and its methods

```
p = subprocess.Popen(  
    args,  
    shell=False,  
    stdin=None, stdout=None, stderr=None,  
    cwd=None,  
    env=None)
```

```
p = subprocess.Popen(  
    args,  
    shell=False,  
    stdin=None, stdout=None, stderr=None,  
    cwd=None,  
    env=None)
```



If `shell=False`, the `Popen` class uses `os.execvp()` to create the new process. In this case, `args` must be a list containing the program name and its arguments

If `shell=True`, `args` is a string that is executed through the shell

```
p = subprocess.Popen(  
    args,  
    shell=False,  
    stdin=None, stdout=None, stderr=None,  
    cwd=None,  
    env=None)
```

If `shell=False`, the `Popen` class uses `os.execvp()` to create the new process. In this case, `args` must be a list containing the program name and its arguments

If `shell=True`, `args` is a string that is executed through the shell

Specify the file handles for standard input, output, and error
Valid values are `PIPE`, an existing file object or descriptor, or `None`.

```
p = subprocess.Popen(  
    args,  
    shell=False,  
    stdin=None, stdout=None, stderr=None,  
    cwd=None,  
    env=None)
```

If `shell=False`, the `Popen` class uses `os.execvp()` to create the new process. In this case, `args` must be a list containing the program name and its arguments

If `shell=True`, `args` is a string that is executed through the shell

Specify the file handles for standard input, output, and error

Valid values are `PIPE`, an existing file object or descriptor, or `None`.

→ If `cwd` is not `None`, the child's current directory will be changed to `cwd` before it is executed


```
p = subprocess.Popen(  
    args,  
    shell=False,  
    stdin=None, stdout=None, stderr=None,  
    cwd=None,  
    env=None)
```

If `shell=False`, the `Popen` class uses `os.execvp()` to create the new process. In this case, `args` must be a list containing the program name and its arguments

If `shell=True`, `args` is a string that is executed through the shell

Specify the file handles for standard input, output, and error
Valid values are `PIPE`, an existing file object or descriptor, or `None`.

If `cwd` is not `None`, the child's current directory will be changed to `cwd` before it is executed

Defines the environment variables for the new process. If `None`, the new process will inherit the current process' environment

The Popen Object

- **Methods:**

- `poll()` - check if the process is still running
- `wait()` - wait until the process terminates
(it may deadlock if PIPEs are full)
- `send_signal(signal)` - send a signal to the process
- `kill()` - kill the process

- **Fields:**

- `stdin, stdout, stderr` – corresponding file objects (only if set to PIPE when the process was created)
- `pid` – the process ID
- `returncode` – contains the process return code (if it is already terminated, or None otherwise)

Example

```
import subprocess

cat = subprocess.Popen(['cat', 'example.py'],
                        stdout=subprocess.PIPE,
                        )

grep = subprocess.Popen(['grep', 'subprocess'],
                        stdin=cat.stdout,
                        stdout=subprocess.PIPE,
                        )

cut = subprocess.Popen(['cut', '-f1', '-d='],
                       stdin=grep.stdout,
                       stdout=subprocess.PIPE,
                       )

end_of_pipe = cut.stdout

for line in end_of_pipe:
    print line.strip()
```

Regular Expression

- The `re` module provides regular expression functionalities
- Regex use `\` to escape special characters: `. ^ $ * + ? { } [] \ | ()`
 - But python use `\` to escape special string sequences: `\n, \t, \r..`
 - Therefore, to match a `\` in a regex you have to write `\\`
 - Or (better) use raw strings: `r'regex'`
- Usual regex syntax:
 - `'^[0-9]+'` : one or more digits at the beginning of the line
- Grouping support:
 - Simple positional group: `(regex)`
 - Named group: `(?P<name>regex)`
 - Match whatever text was matched by the group name: `(?P=name)`

Regex

- If you plan to match the same expression more than once, it's better to compile it to improve performance
- The result of a regex matching is a `MatchObject` instance

```
>>> import re
>>> pattern = re.compile(r'spam:([a-z]+)')
>>> m = pattern.search('this is spam:egg')
>>> m.group(0)      # zero corresponds to the entire regex
spam:egg
>>> m.group(1)      # non-zero is the n-th group in the regex
egg
>>> m.span(1)       # start and end position of a group
(13, 16)
```

Regex

- If you plan to match the same expression more than once, it's better to compile it to improve performance
- The result of a regex matching is a `MatchObject` instance

```
>>> import re
>>> pattern = re.compile(r'spam:(?P<spam>[a-z]+)')
>>> m = pattern.search('this is spam:egg')
>>> m.group(0)      # zero corresponds to the entire regex
spam:egg
>>> m.group("spam")
egg
>>> m.span("spam")  # start and end position of a group
(13, 16)
```

Searching, Matching, and more..

- `match()` determines if the RE matches at the beginning of the string
- `search()` scans through a string, looking for any location where the RE matches
- `finditer()` returns an iterator through all the matches
- `split()` splits the string into a list, splitting it wherever the RE matches
- `findall()` returns all the non-overlapping matches of the pattern as a list of strings
- `sub()` finds all substrings where the RE matches, and replace them with a different string

Regex substitution

- It's possible to substitute a regex with a string

```
>>> pattern = re.compile(r'number: (?P<number>[0-9]+) ')\n>>> pattern.sub('number:XXX', "name:Jack number:06213123")\nname:Jack number:XXX
```


Regex substitution

- It's possible to substitute a regex with a string

```
>>> pattern = re.compile(r'number: (?P<number>[0-9]+) ')
>>> pattern.sub('number:XXX', "name:Jack number:06213123")
name:Jack number:XXX
```

- But it's also possible to substitute with the output of a function

```
def scale(match):
    value = int(match.group('coord'))
    return 'X:%d'%(value*3)

pattern = re.compile(r'X: (?P<coord>[0-9]+) ')
pattern.sub(scale, " X:22 Y:55")
X:66 Y:55
```

More on Regex: Greedy Matching

```
>>> text="<a href='l1'>some text</a>  
      <a href='l2'>more text</a>"  
>>> print re.search("l1'>(.*?)</a>", text).group(1)
```

More on Regex: Greedy Matching


```
>>> text="<a href='l1'>some text</a>  
      <a href='l2'>more text</a>"  
>>> print re.search("l1'>(.*?)</a>", text).group(1)  
some text</a> ... <a href='l2'>more text
```

- By default, the `+` and `*` operators are greedy (i.e., they try to match as many characters as possible)
- Solutions
 - If possible, refine the regular expression (e.g., use `[^<]` instead of `.`)
 - Add a `?` after the operator to make it match as few characters as possible

Lazy Matching

```
>>> text="<a href='l1'>some text</a>  
        <a href='l2'>more text</a>"  
>>> print re.search("l1'>(.*?)</a>", text).group(1)  
some text</a> ... <a href='l2'>more text
```

```
>>> text="<a href='l1'>some text</a>  
        <a href='l2'>more text</a>"  
>>> print re.search("l1'>(.*?)</a>", text).group(1)  
some text
```



Network

- Python standard library includes several modules covering multiple Internet protocols
 - TCP/IP sockets
 - HTTP
 - FTP
 - SSL
 - Telnet
 - XML RPC
 - Mail (IMAP, POP, SMTP)
- External libraries cover the rest:
 - Scapy: powerful low level packet manipulation library
 - Twisted: an event-driven networking engine that supports a large number of protocols (including HTTP, NNTP, DNS, IMAP, SSH, SFTP, IRC, FTP, instant messaging and many others)

Sockets

Server

```
import socket

# create a socket
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)

# associate the socket with a port
s.bind(('', 1234))

# (optional) reuse a socket to prevent waiting..
s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)

# accept connections
s.listen(5)

client_s, addr = s.accept()
print 'Connection from ', addr

client_s.send('Welcome to the server\n')
print client_s.recv(100)
client_s.close()
s.close()
```

Sockets

```
import socket
```

```
# create a socket
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

```
# associate
```

```
s.bind(('localhost', 1234))
```

```
# (optional)
```

```
s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
```

```
# accept
```

```
s.listen(5)
```

```
client_s, client_addr = s.accept()
```

```
print 'Connected from', client_addr
```

```
client_s.send('Welcome to the server!\n')
```

```
print client_s.recv(100)
```

```
client_s.close()
```

```
s.close()
```

```
import socket
```

```
# create a socket
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

```
# open a connection to a certain port
```

```
s.connect(('localhost', 1234))
```

```
print s.recv(100)
```

```
s.send('Bye bye\n')
```

```
s.close()
```

Client

Sockets

```
import socket
```

```
# create a socket
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

```
# associate
```

```
s.bind(('localhost', 1234))
```

```
# (optional)
```

```
s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
```

```
# accept
```

```
s.listen(5)
```

```
client_s, client_addr = s.accept()
```

```
print 'Connected from', client_addr
```

```
client_s.send('Welcome to the server!')
```

```
print client_s.recv(100)
```

```
client_s.close()
```

```
s.close()
```

```
import socket
```

```
# create a socket
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

```
# open a connection to a certain port
```

```
s.connect(('localhost', 1234))
```

```
print s.recv(100)
```

```
s.send('Bye bye\n')
```

```
s.close()
```

Client

Receives up to 100 bytes !!

Sends some characters !!

Use `sendall()` to be sure that
the whole string is sent

Line-based Sockets

- A socket receives and sends sequences of bytes
 - No nice functions like `readline()`
 - Usually implemented manually looping on `recv(1)`
- A way around this problem is to convert the socket to a file-like object

```
s.connect(('localhost', 1234))

# Return a file object associated with the socket
# (for read-write, unbuffered)
fs = s.makefile('rw', 0)

print fs.readline()
fs.write('Bye bye\n')
```

The WEB (client-side)

```
import urllib
import HTMLParser

class GetLinks(HTMLParser.HTMLParser):
    def handle_starttag(self, tag, attrs):
        if tag == 'a':
            for name, value in attrs:
                if name == 'href':
                    print value

gl = GetLinks()

url = 'http://www.iseclab.org/softdev/material.html'
urlconn = urllib.urlopen(url)
urlcontents = urlconn.read()

gl.feed(urlcontents)
```

The WEB (client-side)

```
import urllib
import HTMLParser

class GetLinks(HTMLParser.HTMLParser):
    def handle_starttag(self,tag,attrs):
        if tag == 'a':
            for name,value in attrs:
                if name == 'href':
                    print value

gl = GetLinks()

url = 'http://www.iseclab.org/softdev/material.html'
urlconn = urllib.urlopen(url)
urlcontents = urlconn.read()

gl.feed(urlcontents)
```

The WEB (client-side)

```
import urllib
import HTMLParser

class GetLinks(HTMLParser.HTMLParser):
    def handle_starttag(self, tag, attrs):
        if tag == 'a':
            for name, value in attrs:
                if name == 'href':
                    print value

gl = GetLinks()

url = 'http://www.isecclab.org/softdev/material.html'
urlconn = urllib.urlopen(url)
urlcontents = urlconn.read()

gl.feed(urlcontents)
```

Need to get some data out of some badly-formatted HTML page?
Check out the [Beautiful Soup](#) module

```
opener = urllib2.build_opener()
opener.addheaders = [('User-agent', 'Mozilla/5.0')]
opener.open('http://www.example.com/')
```

Headers

```
import os, cookielib, urllib2
cj = cookielib.MozillaCookieJar()
cj.load(os.path.join(os.environ["HOME"], ".netscape/cookies.txt"))
opener =
urllib2.build_opener(urllib2.HTTPCookieProcessor(cj))
r = opener.open("http://example.com/")
```

Cookies

```
import urllib
import urllib2

url = 'http://www.someserver.com/cgi-bin/register.cgi'
values = {'name' : 'Michael Foord',
          'location' : 'Northampton',
          'language' : 'Python' }

data = urllib.urlencode(values)
req = urllib2.Request(url, data)
response = urllib2.urlopen(req)
the_page = response.read()
```

Post Data

The Web (server-side)

```
import BaseHTTPServer
from SimpleHTTPServer import
SimpleHTTPRequestHandler

HandlerClass.protocol_version = 'HTTP/1.0'
httpd = BaseHTTPServer.HTTPServer(
    ('127.0.0.1', 8080), SimpleHTTPRequestHandler )
httpd.serve_forever()
```

Or simply...

```
> python -m SimpleHTTPServer
```

Python for Webpages (CGI)

```
#!/usr/local/bin/python
import cgi
import cgitb; cgitb.enable()
import Cookie
```

This activates a special exception handler that will display detailed reports in the Web browser if any errors occur

```
print "Content-type: text/html\n"
```

```
form = cgi.FieldStorage()
first_name = form.getvalue('first_name')
last_name = form.getvalue('last_name')
```

Provide access to all the submitted form data (GET and POST)

```
c = Cookie.SimpleCookie(os.environ['HTTP_COOKIE'])
session_id = c['session_id'].value
```

Cookie Management

```
print '''<html><body>
<h1>Title</h1>
...'''
```

Python for Webpages

- Web Frameworks
 - Python data model with object-relational mapper
 - Request routing
 - Automatic administration interface
 - Template system
 - Caching
 - Internationalization
 - <http://wiki.python.org/moin/WebFrameworks> lists 49 web frameworks written in python (Django, Zope, ...)
- Tens of CMS written in python

Threads

- Threads allow a process to do multiple things at once
- A process can have multiple threads
 - Each thread has its own local state
 - All threads share the same global state
- How threads are scheduled to run is dependent on how they are implemented
- Threads in Python are pre-emptive
(i.e, they can be interrupted at any time)
- Two modules:
 - `thread` – primitive thread functionality
 - `threading` - higher-level threading interfaces built on top of the `thread` module

Thread

```
thread.start_new_thread(tfunction , (tparameters))  
lockname = thread.allocate_lock()
```

- Basic threading and locking functionalities
- `start_new_thread()` starts a new python thread
 - The new thread will execute the function `tfunction` invoked with the parameter `tparameters`
 - When the main thread exits, it is system dependent whether the other threads will survive or get killed
- `allocate_lock()` creates a mutex to allows different threads to synchronize

Locking

- `lockname = thread.allocate_lock()`
- `lockname.acquire()` - **takes the lock**
 - Trying to acquire an already taken lock will block the current thread until the lock is released
 - Trying to acquire a lock that was already acquired by the same thread leads to a deadlock situation (!)
- `lockname.release()` - **release the lock**
 - Releasing a lock that was not previously acquired generate a `thread.error`

Example

```
import thread
import socket

main_s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
main_s.bind(('', 1234))
main_s.listen(5)

total = 0
total_lock = thread.allocate_lock()
for x in range(2):
    s, client = main_s.accept()
    thread.start_new_thread(client_manager, (s,))
main_s.close()

print total
```

Example

```
import thread
import socket

main_s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
main_s.bind(('', 1234))
main_s.listen(5)

total = 0
total_lock = thread.allocate_lock()
for x in range(2):
    s, client = main_s.accept()
    thread.start_new_thread(client_manager, (s,))
main_s.close()

print total
```

```
def client_manager(s):
    global total, total_lock
    v = int(s.recv(10))
    total_lock.acquire()
    total += v
    total_lock.release()
    s.close()
```

Example

```
import thread
import socket

main_s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
main_s.bind(('', 1234))
main_s.listen(5)

total = 0
total_lock = thread.allocate_lock()
for x in range(2):
    s, client = main_s.accept()
    thread.start_new_thread(client_manager, (s,))
main_s.close()
```

```
print total
```


When does the
total get printed?

```
def client_manager(s):
    global total, total_lock
    v = int(s.recv(10))
    total_lock.acquire()
    total += v
    total_lock.release()
    s.close()
```

Threading

- Using the threading module, each thread is represented by an instance of the class `Thread`
 - Normally, an application defines a subclass of `Thread` and redefines the method `run()` to implement the thread behavior
- Thread objects
 - `start()` - start the thread
 - `run()` - implement the thread logic (automatically invoked by `start()`)
 - `join([timeout])` – wait till the thread ends (or the optional timeout expires)
 - `is_alive()` - check if the thread is still alive
 - `daemon` – if set to `True` before calling `start()`, set the thread as a daemon thread
- Other useful functions in the threading module:
 - `current_thread()` - returns the current `Thread` object
 - `active_count()` - returns the number of `Thread` object alive
 - `enumerate()` - returns a list of all `Thread` objects currently alive

```
def client_manager(s):  
    global total, total_l  
    v = int(s.recv(10))  
    total_lock.acquire()  
    total += v  
    total_lock.release()  
    s.close()
```



```
class client_manager(threading.Thread):  
    total = 0  
    total_lock = threading.Lock()  
  
    def __init__(self, socket):  
        threading.Thread.__init__(self)  
        self.s = socket  
  
    def run(self):  
        v = int(self.s.recv(10))  
        client_manager.total_lock.acquire()  
        client_manager.total += v  
        client_manager.total_lock.release()  
        self.s.close()
```



```
import thread
import socket

main_s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
main_s.bind(('', 1234))
main_s.listen(5)

total = 0
total_lock = thread.allocate_lock()
for x in range(2):
    s, client = main_s.accept()
    thread.start_new_thread(client_manager, (s,))
main_s.close()

print total
```

```
import threading
import socket

main_s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
main_s.bind(('', 1234))
main_s.listen(5)

for x in range(2):
    s, client = main_s.accept()
    new_client = client_manager(s)
    new_client.start()
main_s.close()
for t in threading.enumerate():
    if t is not threading.current_thread():
        t.join()
print client_manager.total
```

The global interpreter lock

- To facilitate garbage collection, the CPython implementation has a global interpreter lock (GIL) that is used to ensure that only one thread runs at a certain time
- The GIL prevents multiple thread to run in parallel on multi-processor machines
 - It also degrades the performance
(<http://www.dabeaz.com/python/GIL.pdf>)
- So, what if your really need that kind of efficiency?

The global interpreter lock

- To facilitate garbage collection, the CPython implementation has a global interpreter lock (GIL) that is used to ensure that only one thread runs at a certain time
- The GIL prevents multiple thread to run in parallel on multi-processor machines
 - It also degrades the performance
(<http://www.dabeaz.com/python/GIL.pdf>)
- So, what if your really need that kind of efficiency?
 - Maybe you shouldn't use Python in the first place :)
 - You can still use processes instead of threads
 - You can switch to a different Python implementation (e.g., IronPython)

signal, dealing with asynchronous events

- signal handlers can only occur between “atomic” instructions of the Python interpreter
 - signals arriving during long calculations implemented purely in C (such as regular expression matches) may be delayed for an arbitrary amount of time
- Python installs a small number of signal handlers by default:
 - SIGPIPE is ignored (so write errors on pipes and sockets can be reported as ordinary Python exceptions)
 - SIGINT is translated into a KeyboardInterrupt exception
- Signal *cannot* be used to communicate between threads
 - Only the *main thread* can set a new signal handler, and the main thread will be the only one to receive signals

Signals

- `signal.alarm(time)` – a SIGALRM signal will be sent to the process in *time* seconds
- `signal.signal(signalnum, handler)` – set the handler for signal *signalnum* to the function *handler*

```
import signal

# handler for the SIGINT signal (ignore the control-c)
def quit_handler(signum, frame):
    print 'no no no..'

# handler for the alarm
def alarm_handler(signum, frame):
    print 'Wake up'

signal.signal(signal.SIGINT, quit_handler)
signal.signal(signal.SIGALRM, alarm_handler)
signal.alarm(5)

for x in range(10):
    time.sleep(10)
```