

[Software Development]

Python – Optimization and Integration

Davide Balzarotti

Eurecom – Sophia Antipolis, France

When Python is not Enough

- Python is great for rapid application development
- Many famous examples...
 - Fermilab (scientific applications)
 - Google (in many services)
 - Industrial Light & Magic (everything)
 - ...
- But nothing is perfect
 - Sometimes it is just too slow
 - Sometimes libraries are available only in C/C++
 - Sometimes you need to interact with some low-level functionality (system calls, I/O ports..)

The Best of Two Languages

- Two approaches to combine Python with C/C++
- Extending Python
 - Write the application in python and import external plugins to perform functions that are otherwise unavailable or just too slow
 - Best solution when Python is good enough for the rest of the application. The plugins provide the missing pieces
 - Examples: python imaging library (PIL), python numeric

The Best of Two Languages

- Two approaches to combine Python with C/C++
- **Extending Python**
 - Write the application in python and import external plugins to perform functions that are otherwise unavailable or just too slow
 - Best solution when Python is good enough for the rest of the application. The plugins provide the missing pieces
 - Examples: python imaging library (PIL), python numeric
- **Embedding Python**
 - Write the application in C, but run a Python interpreter in it
 - Best approach when Python is not the best solution, but you still want to support configuration or scripting in Python
 - Examples: Civilization IV, Gimp, IDA Pro, GDB

Step I: Find the Bottleneck in the code

- Usually a program spends 90% of its time in 10% of its code
- If the program is too slow, check where the bottleneck is
 - Do not try to guess which is the slow part
(it's not always as easy or as obvious as it may seem)

Step I: Find the Bottleneck in the code

- Usually a program spends 90% of its time in 10% of its code
- If the program is too slow, check where the bottleneck is
 - Do not try to guess which is the slow part
(it's not always as easy or as obvious as it may seem)
- Adopt a more systematic way:
 - Write the entire program prototype in Python
 - Run the program with a profiler (or manually time the different functions)
 - Identify the points that consume most of the CPU time
- Solutions:
 - Check the algorithm and the data structures
(if it is slow, it is often because of a **bad design**.. i.e., it's your fault)
 - If no further optimization is possible... it may be time to get back to C

Timeit

- Simple benchmarking module
 - Useful to measure the performance of small bits of Python code
 - It executes the code multiple times (default is one million) and it takes the total time

```
>>> import timeit
>>> t = timeit.Timer(setup='x=[9,8,7,6,5,4,3,2,1]*100',
                     stmt='x.sort()')
>>> t.timeit(10000)
0.1711

>>> t = timeit.Timer(setup='x=[9,8,7,6,5,4,3,2,1]*100',
                     stmt='y=sorted(x)')
>>> t.timeit(10000)
0.6711
```

Python Profiler

- A profiler is a program that measures the runtime performance of a program, and provides a number of statistics
- Two approaches:
 - **Deterministic**: instrument all the function calls to get precise timing
 - **Statistical**: randomly sample the execution point to estimate where the time is spent in the code

Python Profiler

- A profiler is a program that measures the runtime performance of a program, and provides a number of statistics
- Two approaches:
 - **Deterministic**: instrument all the function calls to get precise timing
 - **Statistical**: randomly sample the execution point to estimate where the time is spent in the code
- Python `cProfile` module: deterministic profiler implemented in C with a reasonable overhead
- Two simple way to use it:
 - From the command line: `python -m cProfile python-script`
 - From the code:

```
import cProfile;
cProfile.run('instruction')
```

```
def distance(p1,p2):  
    x1,y1,z1 = p1  
    x2,y2,z2 = p2  
    tmp = (x2-x1)**2+(y2-y1)**2+(z2-z1)**2  
    return math.sqrt(tmp)  
  
def get_close(points, target, dist):  
    return [p for p in points if distance(target,p) < dist]  
  
points = get_random_points()  
get_close(points, [5,5,5], 20)
```

```
import cProfile
def distance(p1,p2):
    x1,y1,z1 = p1
    x2,y2,z2 = p2
    tmp = (x2-x1)**2+(y2-y1)**2+(z2-z1)**2
    return math.sqrt(tmp)

def get_close(points, target, dist):
    return [p for p in points if distance(target,p) < dist]

points = get_random_points()
cProfile.run('get_close(points, [5,5,5], 20)')
```

```

import cProfile
def distance(p1,p2):
    x1,y1,z1 = p1
    x2,y2,z2 = p2
    tmp = (x2-x1)**2+(y2-y1)**2+(z2-z1)**2
    return math.sqrt(tmp)

def get_close(points, target, dist):
    return [p for p in points if distance(target,p) < dist]

points = get_random_points()
cProfile.run('get_close(points, [5,5,5], 20)')

```

200003 function calls in 0.186 CPU seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.186	0.186	<string>:1(<module>)
100000	0.111	0.000	0.136	0.000	test.py:13(distance)
1	0.050	0.050	0.186	0.186	test.py:20(get_close)
100000	0.024	0.000	0.024	0.000	{math.sqrt}

```

import cProfile
def distance(p1,p2):
    x1,y1,z1 = p1
    x2,y2,z2 = p2
    tmp = (x2-x1)**2+(y2-y1)**2+(z2-z1)**2
    return math.sqrt(tmp)

def get_close(points, target, dist):
    return [p for p in points if distance(target,p) < dist]

points = get_random_points()
cProfile.run('get_close(points, [5,5,5], 20)')

```

200003 function calls in 0.186 CPU seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.186	0.186	<string>:1(<module>)
100000	0.111	0.000	0.136	0.000	test.py:13(distance)
1	0.050	0.050	0.186	0.186	test.py:20(get_close)
100000	0.024	0.000	0.024	0.000	{math.sqrt}

Total time inside the function (excluding the called functions)

```

import cProfile
def distance(p1,p2):
    x1,y1,z1 = p1
    x2,y2,z2 = p2
    tmp = (x2-x1)**2+(y2-y1)**2+(z2-z1)**2
    return math.sqrt(tmp)

def get_close(points, target, dist):
    return [p for p in points if distance(target,p) < dist]

points = get_random_points()
cProfile.run('get_close(points, [5,5,5], 20)')

```

200003 function calls in 0.186 CPU seconds

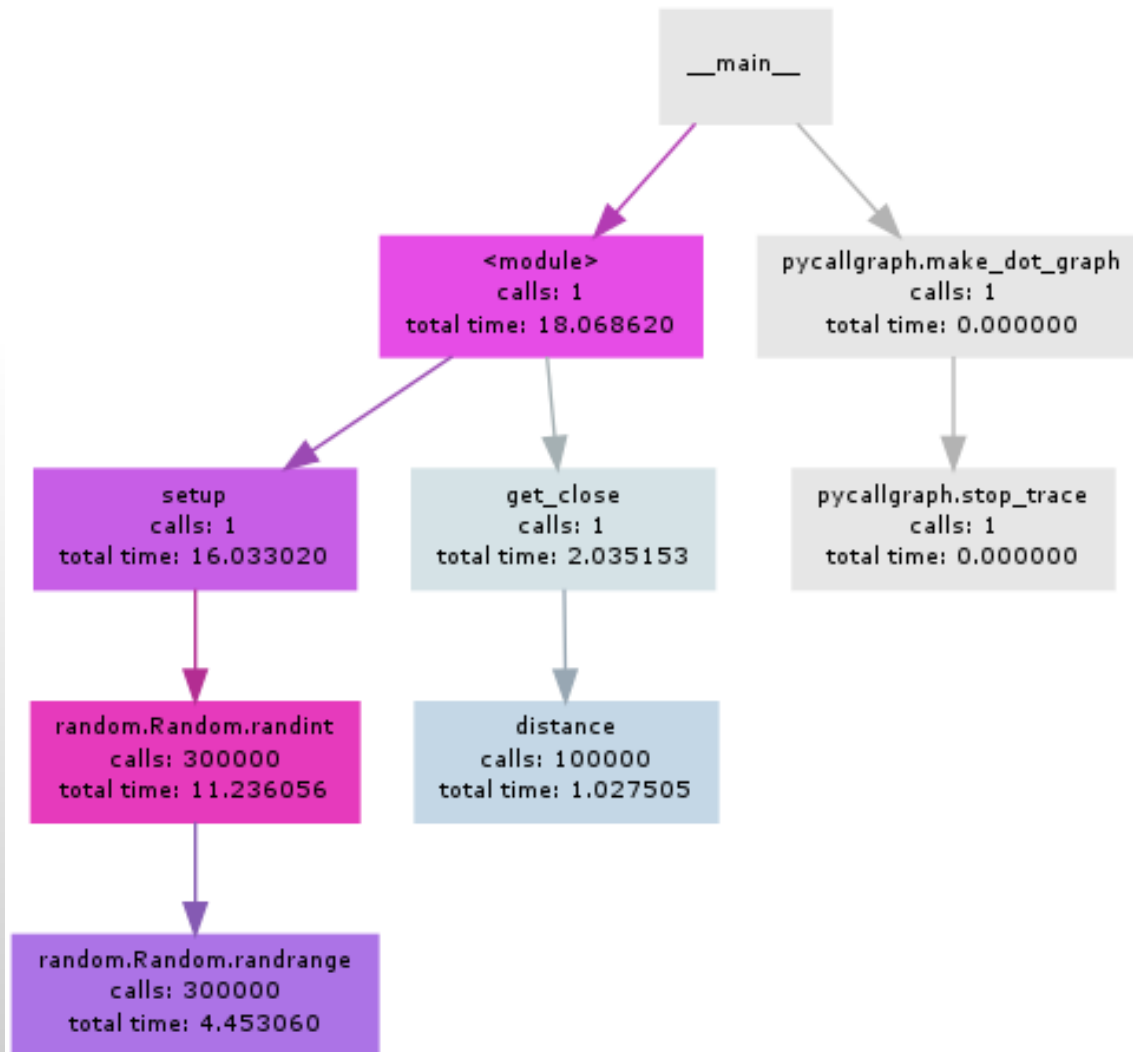
Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.186	0.186	<string>:1(<module>)
100000	0.111	0.000	0.136	0.000	test.py:13(distance)
1	0.050	0.050	0.186	0.186	test.py:20(get_close)
100000	0.024	0.000	0.024	0.000	{math.sqrt}

Total time inside the function (including the called functions)

PyCallGraph

```
> pycallgraph myprogram  
[... a while later...]
```



Python C Extensions

- Most of python builtins and libraries are already implemented in highly optimized C (e.g., `math.sqrt()`) and are quite fast
- Python can be extended by building C/C++ modules that can be imported inside the Python code using the `import` statement
 - Most programs can be greatly optimized in term of speed by rewriting only very small part of them in C
 - Extensions are generally used to provide access to existing C libraries
- A C-coded extension is guaranteed to run only with the version of Python it is compiled for
- A Python extension module named 'foo' generally lives in a dynamic library with the same filename (`foo.so` in Unix)

Distance Module

```
#include "Python.h"

float distance(int x1, int y1, int z1, int x2, int y2, int z2){
    return sqrt((x2-x1)*(x2-x1)+(y2-y1)*(y2-y1)+...));
}

static PyObject
*distance_wrapper (PyObject *self, PyObject *args){
    double xi,yi,zi,xj,yj,zj,dist;
    PyArg_ParseTuple(args,"dddddd",&xi,&yi,&zi,&xj,&yj,&zj);
    dist = distance(xi,yi,zi,xj,yj,zj);
    return Py_BuildValue("d",dist);
}

static PyMethodDef distance_methods[]={
    {"cdistance", distance_wrapper, METH_VARARGS},
    {NULL,NULL}
};

void initdistance(){
    (void) Py_InitModule("distance", distance_methods);
}
```

Distance Module

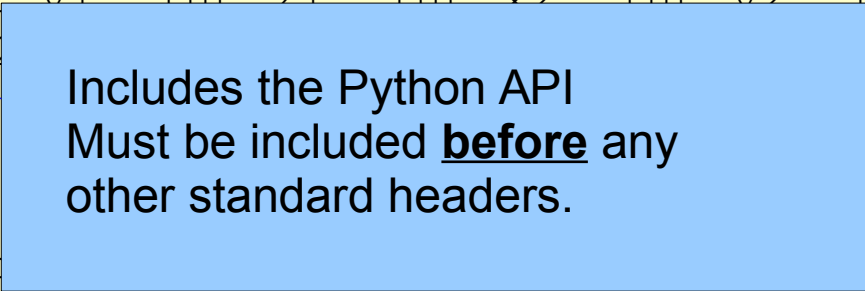
```
#include "Python.h"

float distance(int x1, int y1, int z1, int x2, int y2, int z2) {
    return sqrt((x2-x1)*(x2-x1) + (y2-y1)*(y2-y1) + (z2-z1)*(z2-z1));
}

static PyObject
*distance_wrapper (PyObject *self, PyObject *args) {
    double xi, yi, zi, xj, yj, zj, dist;
    PyArg_ParseTuple(args, "dddddd", &xi, &yi, &zi, &xj, &yj, &zj);
    dist = distance(xi, yi, zi, xj, yj, zj);
    return Py_BuildValue("d", dist);
}

static PyMethodDef distance_methods[]={
    {"cdistance", distance_wrapper, METH_VARARGS},
    {NULL, NULL}
};

void initdistance() {
    (void) Py_InitModule("distance", distance_methods);
}
```



Distance Module

```
#include "Python.h"

float distance(int x1, int y1, int z1, int x2, int y2, int z2){
    return sqrt((x2-x1)*(x2-x1)+(y2-y1)*(y2-y1)+...));
}

static PyObject
*distance_wrapper (PyObject *self, PyObject *args){
    double xi,yi,zi,xj,yj,zj,dist;
    PyArg_ParseTuple(args,"ddddd",&xi,&yi,&zi,&xj,&yj,&zj);
    dist = distance(xi,yi,zi,xj,yj,zj);
    return Py_BuildValue("d",dist);
}

static PyMethodDef distance_methods[]={
    {"cdistance", distance_wrapper, METH_VARARGS},
    {NULL,NULL}
};

void initdistance(){
    (void) Py_InitModule("distance", distance_methods);
}
```

The distance function implemented in C.

Distance Module

```
#include "Python.h"

float distance(int x1, int y1,
              int x2, int y2,
              int z1, int z2){
    return sqrt((x2-x1)*(x2-x1) + (y2-y1)*(y2-y1) + (z2-z1)*(z2-z1));
}
```

Wrapper for the Python code
self is only used in method
args is a tuple containing the arguments

```
static PyObject
*distance_wrapper (PyObject *self, PyObject *args){
    double xi, yi, zi, xj, yj, zj, dist;
    PyArg_ParseTuple(args, "dddddd", &xi, &yi, &zi, &xj, &yj, &zj);
    dist = distance(xi, yi, zi, xj, yj, zj);
    return Py_BuildValue("d", dist);
}
```

PyArg_ParseTuple extract values
from a Python tuple and convert them to
C values (in this case 6 decimals)

Py_BuildValue converts C values to
Python objects. If you have a function
that does not return any value, use the
Py_RETURN_NONE macro

```
static PyMethodDef distance_methods [
    {"distance", distance_wrapper, METH_VARARGS},
    {"__doc__", distance_doc, METH_NOARGS},
    {0, 0, 0}
];

PyInit_distance (void) {
    PyModuleDef_Init(&module);
    module.m_methods = distance_methods;
}
```

Distance Module

```
#include "Python.h"

float distance(int x1, int y1, int z1, int x2, int y2, int z2){
    return sqrt((x1-x2)*(x1-x2)+(y1-y2)*(y1-y2)+(z1-z2)*(z1-z2));
}
```

Method Table:

- python method name
- corresponding function
- calling convention (VARARGS or KEYWORDS)
- docstring

```
static PyObject
*distance_wrapper
    double xi, yi, zi, xj, yj, zj;
    PyArg_ParseTuple(args, "ddd", &xi, &yi, &zi, &xj, &yj, &zj);
    dist = distance(xi, yi, zi, xj, yj, zj);
    return Py_BuildValue("d", dist);
}
```

```
static PyMethodDef distance_methods[]={
    {"cdist", distance_wrapper, METH_VARARGS, 'fast distance'},
    {NULL, NULL, 0, NULL}
};
```

```
void initdistance(){
    (void) Py_InitModule("distance", distance_methods);
}
```

Distance Module

```
#include "Python.h"

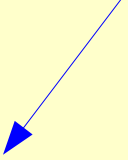
float distance(int x1, int y1, int z1, int x2, int y2, int
z2){
    return sqrt((x2-x1)*(x2-x1)+(y2-y1)*(y2-y1)+...));
}

static PyObject
*distance_wrapper (PyObject *self, PyObject *args){
    double xi,yi,zi,xj,yj,zj,dist;
    PyArg_ParseTuple(args,"dddddd",&xi,&yi,&zi,&xj,&yj,&zj);
    dist = distance(xi,yi,zi,xj,yj,zj);
    return Py_BuildValue("d",dist);
}

static PyMethodDef
{"cdist", dist
{NULL,NULL,0,NULL}
};

void initdistance(){
    (void) Py_InitModule("distance", distance_methods);
}
```

Initialization Function (kind of the main of the module):
Initialize the python module using the method table
(It's the only non-static function defined in the module)



Python Objects

- Most Python/C API functions have one or more arguments as well as a return value of type `PyObject*`
 - `PyObject*` is a pointer to an opaque data type representing an arbitrary Python object
 - `PyObject` cannot be kept on the stack, therefore only pointer variables of type `PyObject*` can be declared
- Memory Management:
 - Python handles memory management automatically by keeping track of all references to variables. This is transparent to the programmer
 - But C is not a memory-managed language !!
As soon as a Python data structure ends up referenced from C, Python lose the ability to track the references automatically

Managing References

- The solution is to have the C code to handle all the reference counters manually
- Two macros:
 - `Py_INCREF ()` - increments the reference counter
 - `Py_DECREF ()` - decrements the reference counter and frees the object when the count reaches zero
- When a piece of code owns a reference to an object, it has to call `Py_DECREF ()` when the reference is no longer needed
 - Forgetting to dispose of an owned reference creates a memory leak

Compiling the Extensions

- `python-config` tells you all you need to use in order to compile a Python module

```
balzarot> gcc `python-config --cflags`  
-shared -fPIC  
-o distance.so distance.c
```

- Now the C module can be used as any other python module

```
balzarot> python  
>>> import distance  
>>> distance.cdist(1,1,1,0,0,0)  
1.732050
```

- By using the `cdist` function, the `close()` function is now **100%** faster

The Pythonic Way: Distutils

- The preferred approach for building a Python extension module is to compile it with distutils
 - Distutils takes care of building the extension with all the correct flags, headers, ...
 - It can also handle the installation of the package
- The developer has to write a `setup.py` file

```
from distutils.core import setup, Extension

module1 = Extension('distance', sources=['distance.c'])

setup (name = 'distance',
       version = '1.0',
       description = 'This is the distance package',
       ext_modules = [module1])
```

The Pythonic Way: Distutils

- The preferred approach for building a Python extension module is to compile it with distutils
 - Distutils takes care of building the extension with all the correct flags, headers, ...
 - It can also handle the installation of the package
- The developer has to write a `setup.py` file

```
from distutils.core import setup, Extension  
  
module1 = Extension('distance', sources=['distance.c'])  
  
setup (name = 'distance',
```

```
balzarot> python setup.py build_ext --inplace  
balzarot> python setup.py install
```

Using Existing Libraries

- Manual functions decoration can be cumbersome
 - It is appropriate when coding new built-in data types or core Python extensions...
 - ... but if you have a library with hundreds of functions written in C and you want to import it in Python, it's better to use a tool to simplify the job
- Examples:
 - OpenGL module
 - Subversion
 - Mapserver
 - Fifengine
 - ...

SWIG

- The *Simplified Wrapper and Interface Generator* (SWIG) is an interface compiler that connects programs written in C/C++ with scripting languages such as Perl, Python, Ruby, and Tcl (<http://www.swig.org>)
- SWIG decorates C source with the necessary Python markup
- Generates two files:
 - `extension.py` Python file that contains high-level support code. This is the file that you will import to use the module
 - `extension_wrap.c` C source file containing the low-level wrappers that need to be compiled and linked with the rest of the C library to obtain the Python module
- Limitations:
 - C++ support is not perfect.. check out Boost .Python for C++

SWIG

distance.c

```
#import "distance.h"  
#import <math.h>  
  
float distance(int x1, int y1, int z1, int x2, int y2, int  
z2){  
    return sqrt((x2-x1)*(x2-x1)+(y2-y1)*(y2-y1)+...));  
}
```

distance.h

```
float distance(int x1, int y1, int z1, int x2, int y2, int z2);
```

SWIG

distance.c

```
#import "distance.h"
#import <math.h>

float distance(int x1, int y1, int z1, int x2, int y2, int
z2){
    return sqrt((x2-x1)*(x2-x1)+(y2-y1)*(y2-y1)+...));
}
```

distance.h

```
float distance(int x1, int y1, int z1, int x2, int y2, int z2);
```

distance.i

```
%module distance

%{
#define SWIG_FILE_WITH_INIT
#include "distance.h"
%}

extern float distance(int x1, int y1, int z1, int x2, int
y2, int z2);
```

SWIG – Building the Module

```
balzarot> swig -python distance.i
balzarot> ls
distance.c      distance.h      distance.py
distance_wrap.c distance.i
```

→ Run swig to generate Python wrappers

SWIG – Building the Module

```
balzarot> swig -python distance.i
balzarot> ls
distance.c      distance.h  distance.py
distance_wrap.c distance.i

balzarot> gcc distance.c distance_wrap.c
`python-config --cflags` -shared -fPIC
-o _distance.so
```

Compile the library and the Python wrapper

→ **Important:** the name of the output file has to match the name of the module **prefixed by an underscore**

SWIG – Building the Module

```
balzarot> swig -python distance.i
balzarot> ls
distance.c      distance.h  distance.py
distance_wrap.c distance.i

balzarot> gcc distance.c distance_wrap.c
`python-config --cflags` -shared -fPIC
-o _distance.so

balzarot> python
>>> import distance
>>> distance.distance
<built-in function distance>
```

→ Use the module.. the usual way

SWIG – Global Variables

- There is no direct way to map variable assignment in C to variable assignment in Python
 - In Python: "x=3; y=x" x and y are names for the same object containing the value 3
 - In C: "x=3; y=x;" x and y refer to different memory locations that store the value 3
- SWIG can generate wrappers to access global C variable in Python as attributes of a special object called cvar (this name can be changed using the -globals option)

```
%module example
%inline %{
extern int x;
%}
```

```
>>> import example
>>> example.cvar.x = 3
```

SWIG – Structures

- SWIG automatically wraps C structures in Python classes

```
%module example  
struct Vector {  
    double x,y,z;  
};
```

```
>>> import example  
>>> v = example.Vector()  
>>> v.x = 2.3  
>>> v.y = 3.1
```

PyPy

- PyPy is a fast, compliant alternative implementation of the Python language (2.7.10 and 3.2.5 beta)
- Advantages:
 - Increases speed thanks to its Just-in-Time compilation
 - Slightly reduced memory usage due to optimized GC
- Check features and compatibility on <http://pypy.org>
- Very easy to use: just run `pypy` instead of `python`

Embedding Python

- Include Python code in a C/C++ application
- Very useful for:
 - User defined routine
 - Extension scripts
- And it is much simpler than you may think

Embedding Python

- Include Python code in a C/C++ application
- Very useful for:
 - User defined routine
 - Extension scripts
- And it is much simpler than you may think

```
#include "Python.h"

void hello_python()
{
    Py_Initialize();
    PyRun_SimpleString("print 'Hello World'");
    Py_Finalize();
}

void main() {
    hello_python();
}
```

A (slightly) More Sophisticated Example

- Executing Python code is useful when it can interact with the rest of the application
- For example:
 - A Python script that modify an image in Gimp
 - A Python script that implement the artificial intelligence of a monster in a game
 - ...
- This can be done by combining what we already saw about Python extensions with an embedded interpreter


```

#include "Python.h"

static PyObject*
emb_hello(PyObject *self, PyObject *args)
{
    printf("Hello From C\n");
    return Py_BuildValue("i", 1);
}

static PyMethodDef EmbMethods[] = {
    {"emb_hello", emb_hello, METH_VARARGS, ""},
    {NULL, NULL, 0, NULL}
};

void hello_python()
{
    FILE* f;
    f = fopen("script.py", "r");
    Py_Initialize();
    Py_InitModule("emb", EmbMethods);
    PyRun_SimpleFile(f, "script.py");
    Py_Finalize();
}

void main() {
    hello_python();
}

```

Python
extension to
allow the script
to interact with
the rest of the
application

Embed the
Python
interpreter to
execute external
scripts

```

#include "Python.h"

static PyObject*
emb_hello(PyObject *self, PyObject *args)
{
    printf("Hello From C\n");
    return Py_BuildValue("i", 1);
}

static PyMethodDef EmbMethods[] = {
    {"emb_hello", emb_hello, METH_VARARGS, ""},
    {NULL, NULL, 0, NULL}
};

void hello_python()
{
    FILE* f;
    f = fopen("script.py", "r");
    Py_Initialize();
    Py_InitModule("emb", EmbMethods);
    PyRun_SimpleFile(f, "script.py");
    Py_Finalize();
}

void main() {
    hello_python();
}

```

```
#script.py
```

```
import emb;
```

```
print 'Hello From Python'
```

```
emb.emb_hello()
```